

# Project

First, we have an image with width=850 and height=567:



Figure 1: A simple image

## 1 CPU implementation

### 1.1 Logic steps

There were five steps that I did in this project:

- Pad the image based on the window's size.

- Convert the image from RGB to HSV.
- For each pixel, extract its corresponding four windows.
- Calculate the standard deviation of each window based on the V channels using `numpy.std`.
- Assign R, G, B values for each pixel based on the mean value of the windows (`numpy.mean`) with the smallest standard deviation.

## 1.2 Implementation steps

### 1.2.1 Pad the image based on the window's size

```

1 h_pad = args.window_size // 2
2 w_pad = args.window_size // 2
3 image = np.pad(
4     image,
5     pad_width=(
6         (h_pad, h_pad),
7         (w_pad, w_pad),
8         (0, 0),
9     ),
10    mode="constant",
11    constant_values=0,
12 )

```

### 1.2.2 Convert the image from RGB to HSV

```

1 def rgb_to_hsv(image):
2     y, x = image.shape[0], image.shape[1]
3     output = np.zeros((image.shape), np.float32)
4     for i in range(y):
5         for j in range(x):
6             max_value = max(image[i, j, 0], image[i, j, 1], image[i, j,
7             2])
8             min_value = min(image[i, j, 0], image[i, j, 1], image[i, j,
9             2])
10            delta = max_value - min_value

```

```

9         if delta == 0:
10             h_value = 0
11         elif max_value == image[i, j, 0]:
12             h_value = 60 * (((image[i, j, 1] - image[i, j, 2]) /
delta) % 6)
13         elif max_value == image[i, j, 1]:
14             h_value = 60 * (((image[i, j, 2] - image[i, j, 0]) /
delta) + 2)
15         elif max_value == image[i, j, 2]:
16             h_value = 60 * (((image[i, j, 0] - image[i, j, 1]) /
delta) + 4)
17
18         if max_value == 0:
19             s_value = 0
20         else:
21             s_value = delta / max_value
22         v_value = max_value
23         output[i, j, 0] = h_value
24         output[i, j, 1] = s_value
25         output[i, j, 2] = v_value
26     return output

```

### 1.2.3 Extract four corresponding windows

```

1 def extract_window(img, top, height, left, width):
2     return img[top : top + height, left : left + width]
3
4 window_coordinates = [
5     [
6         h - small_window_height + 1,
7         small_window_height,
8         w - small_window_width + 1,
9         small_window_width,
10    ],
11    [
12        h - small_window_height + 1,
13        small_window_height,
14        w,
15        small_window_width,

```

```

16 ],
17 [
18     h,
19     small_window_height,
20     w - small_window_width + 1,
21     small_window_width,
22 ],
23 [h, small_window_height, w, small_window_width],
24 ]
25
26 window1 = extract_window(image_v, *window_coordinates[0])
27 window2 = extract_window(image_v, *window_coordinates[1])
28 window3 = extract_window(image_v, *window_coordinates[2])
29 window4 = extract_window(image_v, *window_coordinates[3])

```

#### 1.2.4 Calculate the standard deviation of each window

```

1 std_dev1 = np.std(window1)
2 std_dev2 = np.std(window2)
3 std_dev3 = np.std(window3)
4 std_dev4 = np.std(window4)
5 min_std = min(std_dev1, std_dev2, std_dev3, std_dev4)

```

#### 1.2.5 Assign the corresponding *R, G, and B* values for each pixel

```

1 if std_dev1 == min_std:
2     mean_r = np.mean(extract_window(image_r, *window_coordinates[0]))
3     mean_g = np.mean(extract_window(image_g, *window_coordinates[0]))
4     mean_b = np.mean(extract_window(image_b, *window_coordinates[0]))
5 elif std_dev2 == min_std:
6     mean_r = np.mean(extract_window(image_r, *window_coordinates[1]))
7     mean_g = np.mean(extract_window(image_g, *window_coordinates[1]))
8     mean_b = np.mean(extract_window(image_b, *window_coordinates[1]))
9 elif std_dev3 == min_std:
10    mean_r = np.mean(extract_window(image_r, *window_coordinates[2]))
11    mean_g = np.mean(extract_window(image_g, *window_coordinates[2]))
12    mean_b = np.mean(extract_window(image_b, *window_coordinates[2]))
13 elif std_dev4 == min_std:
14    mean_r = np.mean(extract_window(image_r, *window_coordinates[3]))

```

```

15     mean_g = np.mean(extract_window(image_g, *window_coordinates[3]))
16     mean_b = np.mean(extract_window(image_b, *window_coordinates[3]))
17 image_output[
18     h - small_window_height + 1, w - small_window_height + 1, 0
19 ] = mean_r
20 image_output[
21     h - small_window_height + 1, w - small_window_height + 1, 1
22 ] = mean_g
23 image_output[
24     h - small_window_height + 1, w - small_window_height + 1, 2
25 ] = mean_b

```

### 1.2.6 Result

We have the resulting image:




Figure 2: Applying Kuwahara filter on the image using CPU

## 2 GPU implementation (without shared memory)

### 2.1 Logic steps

There were five steps that I did in this project:

- Pad the image based on the window's size.
- Convert the image from RGB to HSV.
- For each pixel, extract its corresponding four windows.
- Calculate the standard deviation of each window based on the V channels (from scratch).
- Assign R, G, and B values for each pixel based on the mean value of the windows (from scratch) that has the smallest standard deviation.

### 2.2 Implementation steps

#### 2.2.1 Pad the image based on the window's size

This step is the same as in the CPU

#### 2.2.2 Convert the image from RGB to HSV

```
1 @cuda.jit
2 def rgb_to_hsv(src, dst):
3     x, y = cuda.grid(2)
4     if y < dst.shape[0] and x < dst.shape[1]:
5         max_value = max(src[y, x, 0], src[y, x, 1], src[y, x, 2])
6         min_value = min(src[y, x, 0], src[y, x, 1], src[y, x, 2])
7         delta = max_value - min_value
8         if delta == 0:
9             h_value = 0
10        elif max_value == src[y, x, 0]:
11            h_value = 60 * (((src[y, x, 1] - src[y, x, 2]) / delta) % 6)
12        elif max_value == src[y, x, 1]:
13            h_value = 60 * (((src[y, x, 2] - src[y, x, 0]) / delta) + 2)
```

```

14         elif max_value == src[y, x, 2]:
15             h_value = 60 * (((src[y, x, 0] - src[y, x, 1]) / delta) + 4)
16
17         if max_value == 0:
18             s_value = 0
19         else:
20             s_value = delta / max_value
21         v_value = max_value
22         dst[y, x, 0] = h_value
23         dst[y, x, 1] = s_value
24         dst[y, x, 2] = v_value

```

### 2.2.3 Extract four corresponding windows

This step is a bit different from the CPU version. Instead of having four windows with specific sizes, I just calculate their coordinates(top, left, width, and height)

```

1 tops = cuda.local.array(4, numba.int64)
2 heights = cuda.local.array(4, numba.int64)
3 lefts = cuda.local.array(4, numba.int64)
4 widths = cuda.local.array(4, numba.int64)
5
6 tops[0] = y - small_window_height + 1
7 tops[1] = y - small_window_height + 1
8 tops[2] = y
9 tops[3] = y
10
11 heights[0] = small_window_height
12 heights[1] = small_window_height
13 heights[2] = small_window_height
14 heights[3] = small_window_height
15
16 lefts[0] = x - small_window_width + 1
17 lefts[1] = x
18 lefts[2] = x - small_window_width + 1
19 lefts[3] = x
20
21 widths[0] = small_window_width
22 widths[1] = small_window_width
23 widths[2] = small_window_width

```



```
24 widths[3] = small_window_width
```

#### 2.2.4 Calculate the standard deviation of each window and get the coordinate of the window that has the smallest deviation

```
1 smallest_std_window = np.inf
2 smallest_window_idx = -1
3 for window in range(4):
4     total_sum_window = 0
5     top = tops[window]
6     left = lefts[window]
7     height = heights[window]
8     width = widths[window]
9     for i in range(top, top + height):
10         for j in range(left, left + width):
11             total_sum_window += image_v[i, j]
12
13     mean_window = total_sum_window / (width * height)
14     sum_of_squared_diff_window = 0
15
16     for i in range(top, top + height):
17         for j in range(left, left + width):
18             diff = image_v[i, j] - mean_window
19             sum_of_squared_diff_window += diff * diff
20
21     std_window = math.sqrt(sum_of_squared_diff_window / (width * height)
22                             )
23     if std_window < smallest_std_window:
24         smallest_std_window = std_window
25         smallest_window_idx = window
26
27 top = tops[smallest_window_idx]
28 left = lefts[smallest_window_idx]
29 height = heights[smallest_window_idx]
30 width = widths[smallest_window_idx]
```

#### 2.2.5 Assign the corresponding R, G, and B values for each pixel



```

1 total_sum_window_r = 0.0
2 total_sum_window_g = 0.0
3 total_sum_window_b = 0.0
4
5 for i in range(top, top + 4):
6     for j in range(left, left + 4):
7         total_sum_window_r += src_rgb[i, j, 0]
8         total_sum_window_g += src_rgb[i, j, 1]
9         total_sum_window_b += src_rgb[i, j, 2]
10
11 dst[y, x, 0] = total_sum_window_r / (width * height)
12 dst[y, x, 1] = total_sum_window_g / (width * height)
13 dst[y, x, 2] = total_sum_window_b / (width * height)

```

We have the resulting image:



Figure 3: Applying Kuwahara filter on the image using GPU (without shared memory)

## 3 GPU implementation (with shared memory)

### 3.1 Logic steps

There were five steps that I did in this project:

- Pad the image based on the window's size.
- Convert the image from RGB to HSV.
- For each pixel, extract its corresponding four windows.
- Calculate the standard deviation of each window based on the V channels (from scratch).
- Assign R, G, and B values for each pixel based on the mean value of the windows (from scratch) with the smallest standard deviation.

### 3.2 Implementation steps

#### 3.2.1 Pad the image based on the window's size

I pad the image so that the image's width and height are divisible by 8.

```
1 def pad_image_to_divisible_by_8(image, n):
2     padded_height = image.shape[0] + 2 * n
3     padded_width = image.shape[1] + 2 * n
4
5     left_pad = top_pad = n
6     right_pad = bottom_pad = n
7
8     if padded_width % 8 != 0:
9         additional_width = 8 - (padded_width % 8)
10        left_pad += additional_width // 2
11        right_pad += additional_width - (additional_width // 2)
12
13    if padded_height % 8 != 0:
14        additional_height = 8 - (padded_height % 8)
15        top_pad += additional_height // 2
16        bottom_pad += additional_height - (additional_height // 2)
```

```

17
18     padded_image = np.pad(
19         image, ((top_pad, bottom_pad), (left_pad, right_pad), (0, 0)),
20         mode="constant"
21     )
22     return padded_image, left_pad, right_pad, top_pad, bottom_pad
23 image, left_pad, right_pad, top_pad, bottom_pad =
24     pad_image_to_divisible_by_8(
25         image, h_pad

```

### 3.2.2 Convert the image from RGB to HSV

```

1 x, y = cuda.grid(2)
2 tx = cuda.threadIdx.x
3 ty = cuda.threadIdx.y
4 shared_hsv = cuda.shared.array(shape=(8, 8, 3), dtype=numba.float32)
5
6 if y < dst.shape[0] and x < dst.shape[1]:
7     for i in range(3):
8         shared_hsv[ty, tx, i] = src[y, x, i]
9 cuda.syncthreads()
10 if ty < 8 and tx < 8:
11     max_value = max(
12         shared_hsv[ty, tx, 0], shared_hsv[ty, tx, 1], shared_hsv[ty, tx,
13         2]
14     )
15     min_value = min(
16         shared_hsv[ty, tx, 0], shared_hsv[ty, tx, 1], shared_hsv[ty, tx,
17         2]
18     )
19     delta = max_value - min_value
20     if delta == 0:
21         h_value = 0
22     elif max_value == shared_hsv[ty, tx, 0]:
23         h_value = 60 * (
24             ((shared_hsv[ty, tx, 1] - shared_hsv[ty, tx, 2]) / delta) %
25
6

```

```

23         )
24     elif max_value == shared_hsv[ty, tx, 1]:
25         h_value = 60 * (
26             ((shared_hsv[ty, tx, 2] - shared_hsv[ty, tx, 0]) / delta) +
27             2
28         )
29     elif max_value == shared_hsv[ty, tx, 2]:
30         h_value = 60 * (
31             ((shared_hsv[ty, tx, 0] - shared_hsv[ty, tx, 1]) / delta) +
32             4
33         )
34     if max_value == 0:
35         s_value = 0
36     else:
37         s_value = delta / max_value
38     v_value = max_value
39     dst[y, x, 0] = h_value
40     dst[y, x, 1] = s_value
41     dst[y, x, 2] = v_value
42 cuda.syncthreads()

```

### 3.2.3 Extract four corresponding windows

This step is a bit different from the CPU version. Instead of having four windows with specific sizes, I just calculate the coordinates of them (top, left, width, and height)

```

1 tx = cuda.threadIdx.x
2 ty = cuda.threadIdx.y
3 shared_hsv = cuda.shared.array(shape=(26, 26, 1), dtype=numba.float32)
4 shared_rgb = cuda.shared.array(shape=(26, 26, 3), dtype=numba.float32)
5 if (
6     x < src_hsv.shape[1]
7     and y < src_hsv.shape[0]
8 ):
9     shared_hsv[ty, tx] = src_hsv[
10         y - 3,
11         x - 3,
12     ]

```

```

13     shared_hsv[ty + 6, tx] = src_hsv[
14         y + 3,
15         x - 3,
16     ]
17     shared_hsv[ty, tx + 6] = src_hsv[
18         y - 3,
19         x + 3,
20     ]
21     shared_hsv[ty + 6, tx + 6] = src_hsv[
22         y + 3,
23         x + 3,
24     ]
25
26 if (
27     x < src_rgb.shape[1]
28     and y < src_rgb.shape[0]
29 ):
30     for i in range(3):
31         shared_rgb[ty, tx, i] = src_rgb[y - 3, x - 3, i]
32         shared_rgb[ty + 6, tx, i] = src_rgb[y + 3, x - 3, i]
33         shared_rgb[ty, tx + 6, i] = src_rgb[y - 3, x + 3, i]
34         shared_rgb[ty + 6, tx + 6, i] = src_rgb[y + 3, x + 3, i]
35
36 cuda.syncthreads()
37 if (
38     x < src_hsv.shape[1]
39     and y < src_hsv.shape[0]
40 ):
41     tops = cuda.local.array(4, numba.int64)
42     heights = cuda.local.array(4, numba.int64)
43     lefts = cuda.local.array(4, numba.int64)
44     widths = cuda.local.array(4, numba.int64)
45
46     tops[0] = ty
47     tops[1] = ty
48     tops[2] = ty + small_window_height - 1
49     tops[3] = ty + small_window_height - 1
50
51     heights[0] = small_window_height
52     heights[1] = small_window_height

```

```

53     heights[2] = small_window_height
54     heights[3] = small_window_height
55
56     lefts[0] = tx
57     lefts[1] = tx + small_window_width - 1
58     lefts[2] = tx
59     lefts[3] = tx + small_window_width - 1
60
61     widths[0] = small_window_width
62     widths[1] = small_window_width
63     widths[2] = small_window_width
64     widths[3] = small_window_width

```

### 3.2.4 Calculate the standard deviation of each window and get the coordinate of the window that has the smallest deviation

```

1  smallest_std_window = np.inf
2  smallest_window_idx = -1
3
4  for window in range(4):
5      total_sum_window = np.float32(0)
6      top = tops[window]
7      left = lefts[window]
8      height = heights[window]
9      width = widths[window]
10     for i in range(top, top + height):
11         for j in range(left, left + width):
12             total_sum_window += shared_hsv[i, j, 0]
13     mean_window = total_sum_window / (width * height)
14     sum_of_squared_diff_window = 0
15
16     for i in range(top, top + height):
17         for j in range(left, left + width):
18             diff = shared_hsv[i, j, 0] - mean_window
19             sum_of_squared_diff_window += diff * diff
20
21     std_window = math.sqrt(sum_of_squared_diff_window / (width * height)
22 )
23     if std_window < smallest_std_window:

```

```

23         smallest_std_window = std_window
24         smallest_window_idx = window
25
26 top = tops[smallest_window_idx]
27 left = lefts[smallest_window_idx]
28 height = heights[smallest_window_idx]
29 width = widths[smallest_window_idx]

```

### 3.2.5 Assign the corresponding R, G, and B values for each pixel

```

1 total_sum_window_r = np.float32(0)
2 total_sum_window_g = np.float32(0)
3 total_sum_window_b = np.float32(0)
4
5 for i in range(top, top + 4):
6     for j in range(left, left + 4):
7         total_sum_window_r += shared_rgb[i, j, 0]
8         total_sum_window_g += shared_rgb[i, j, 1]
9         total_sum_window_b += shared_rgb[i, j, 2]
10 dst[y, x, 0] = total_sum_window_r / (width * height)
11 dst[y, x, 1] = total_sum_window_g / (width * height)
12 dst[y, x, 2] = total_sum_window_b / (width * height)

```

We have the resulting image:





Figure 4: Applying Kuwahara filter on the image using GPU (with shared memory)

## 4 Run time comparison

All experiments are run on Intel(R) Core(TM) i7-9750H CPU @ 2.60GHz and GTX 1650.

	Ws 3	Ws 5	Ws 7	Ws 9
CPU	53.12	56.35	58.46	51.58
GPU (without shared memory)	0.6169	0.5848	0.6291	0.6349
GPU (with shared memory)	0.5844	0.6296	0.5837	0.6678

Table 1: Runtime comparison between 3 different methods on four different window sizes (Ws: window size).

I compare the run time on CPU, GPU (without shared memory), and GPU (with shared memory) with five different window sizes. For the sake of simplicity, I only

compare the runtime of Kuwahara filter and not other functions. The execution of the filter on the CPU is very slow, around 51-60s for different window sizes. I noticed an interesting thing: Large windows size do not always run slower than small windows size. I decrease the run time by nearly 90-100 times when I run the code on GPU (without shared memory). Using shared memory results in nearly the same run times as not using shared memory (same block size = 8).

## 5 Conclusion and future work

Things that I have done in this project:

- Implementation of Kuwahara filter on CPU, GPU (without shared memory), and GPU (with shared memory).
- No hardcoded of most values (except block size and shared array's shape).

Things that I have not done in this project:

- Optimization of Kuwahara filter on GPU.